

62%

wins are ensemble stacking

35/45

GBM the dominant model

$\kappa = 0.65$

coding reliability

45%

winner–control author overlap

5 of 6

folklore couplings fail

Abstract

Does elaborate feature engineering win tabular Kaggle competitions? We coded 45 winning solutions (2022–2026) and 11 near-winning controls across a 53-technique taxonomy. Three findings:

- Winning reduces to a few **paradigms** — ensemble stacking dominates (28/45).
- Winning feature engineering is a **small shared core** — it does *not cleanly* separate winners from near-winners ($n = 11$, n.s.).
- The documented top is a small recurring **guild**.

Winning turns less on FE *volume* than on paradigm fit and community membership.

Keywords: *meta-analysis, Kaggle, tabular ML, feature engineering, ensembling, content analysis*

Introduction & Motivation

“What wins on Kaggle” is mostly **anecdotal**.

Benchmarks compare *algorithms* on fixed data — never winners against strong non-winners. We code both.

Research Questions

- RQ1** — What strategy *paradigms* define winning, and do claimed constraint→strategy rules hold?
- RQ2** — Which FE techniques do winners use, and how much?
- RQ3** — Does FE *distinguish* winners from strong non-winners?
- RQ4** — How is the winning population *structured*?

Data Description

Attribute	Description
Source	Kaggle (Playground S3–S6, TPS 2022, 1 Featured)
Sample size	45 winners + 11 near-winning controls
Variables	40-column schema + 53-technique FE taxonomy
Time range	Feb 2022 – Apr 2026
Selection	voluntary writeups ⇒ documentation-self-selection bias

Exploratory Data Analysis

Strategies are not evenly spread: **ensemble stacking** takes most wins, and gradient-boosted trees are the dominant model in **35 of 45** entries.




Figure 1. Paradigm distribution across the 45 winning solutions.

Methodology

Three-pass content-analysis coding of each solution, with a near-winner *control group* and a *blind reliability re-code*.




Figure 2. Multi-pass coding workflow with a near-winner control group and reliability check.

Analytical Framework

No model is trained — the unit is one completed solution (no RMSE/accuracy). The framework:

- Four paradigms** — ensemble-stacking, single-model-FE, lookup-exploit, problem-fit-NN.
- Winner–control design** — paired FE comparison (sign test).
- Coupling audit** — six folklore “constraint→strategy” claims tested.

Each winner is tested against a **near-winning control** from the *same* competition — the comparison that isolates what distinguishes *winning* from merely strong play.

Implementation & Tools

Component	Tools / Libraries
Language / env	Python 3.13, Poetry
Analysis	pandas, numpy, Python standard library
Collection	Kaggle API, requests + BeautifulSoup
Coding	anthropic SDK + Claude Code (AI-assisted, researcher-directed)
Visualization	matplotlib

Results

FE does not separate winners from near-winners ($n = 11$, not significant); only *learned* features lean winner.




Figure 3. FE-family prevalence, winners vs. near-winning controls.

5 of 6 folklore couplings fail — only linear-stacker dominance holds.




Figure 4. Coupling-evidence verdicts across the six promoted couplings.

Discussion & Conclusions

Key Takeaways

- Winning $\neq$ FE folklore — a small shared core, not elaborate craft.
- Learned features** — the one edge that leans to winners.
- The top is a **recurring guild** (45% author overlap) — partly about *who competes*.

Limitations: small n ; synthetic-data corpus; intra-rater reliability; documentation-self-selection.

Future Work: real-data Featured competitions; other platforms and modalities.

References

- Caruana, R., & Niculescu-Mizil, A. (2006). An empirical comparison of supervised learning algorithms. *ICML*, 161–168.
- Gorishniy, Y., et al. (2021). Revisiting deep learning models for tabular data. *NeurIPS* 34.
- Borisov, V., et al. (2024). Deep neural networks and tabular data: A survey. *IEEE TNLS*, 35(6).
- Makridakis, S., Spiliotis, E., & Assimakopoulos, V. (2022). M5 accuracy competition: Results, findings, conclusions. *Int. J. Forecasting*, 38(4).
- Merton, R. K. (1968). The Matthew effect in science. *Science*, 159(3810).
- Zhang, T., et al. (2023). OpenFE: Automated feature generation with expert-level performance. *ICML*.

Acknowledgments & Contact

The author thanks **Dr. Pedro Albuquerque** for his guidance throughout CS 495, and the Bellevue College Computer Science program for supporting this capstone.

Contact: kenneth.young@bellevuecollege.edu

Code / Repository: [\[github.com/...\]](#) — *add URL*